


Now Assist AI Agents Prompting Guide



Yokohama and Zurich release

Author: Victor Chen

www.servicenow.com

© 2025 ServiceNow, Inc. All rights reserved. ServiceNow, the ServiceNow logo, and other ServiceNow marks are trademarks and/or registered trademarks of ServiceNow, Inc. in the United States and/or other countries. Other company names, product names, and logos may be trademarks of the respective companies with which they are associated.


Table of contents

Introduction.....	4
What are AI agents	4
How to use this guide	4
Acknowledgements.....	5
1. How do Now Assist AI agents work	7
Whether to build one large AI agent or several specific AI agents.....	9
2. General prompting guidance	10
Getting prepared	10
Structuring and formatting your prompts	10
Prompting for agentic workflows	12
Prompting for AI agents	12
Other prompting techniques.....	15
Referencing ServiceNow data	15
Handing over from one agent to another	15
Using memory.....	16
End-user inputs and outputs	16
Using Now Assist to create a prompt.....	18
3. Prompting with Tools.....	20
List of internal AI agent tools	22
List of keyword verbs to use for tools	22
4. More prompt examples	23
5. Prompting troubleshooting	29
6. Additional Resources	31

Introduction

What are AI agents

AI agents are “digital workers” that enhance your organization's productivity by intelligently orchestrating workflows, skills, and actions across your business.



ServiceNow agentic workflows and AI agents are powered by prompting and large language models (LLM). Prompts are the instructions given to the agentic workflows and AI agents on how to orchestrate as a team to solve business problems or achieve objectives set by a developer.

What is a prompt? According to OpenAI:

A prompt for a Large Language Model (LLM) is a text input that initiates a conversation or triggers a response from the model. Prompt engineering is the process of designing and optimizing input prompts to effectively guide a language model's responses.

How to use this guide

The good news is that you don't need to be an AI engineer or data scientist to be a good prompt writer. Prompt writing is an iterative and sometimes experimental process. As you are building and deploying your own agentic workflows and AI agents on the ServiceNow AI Platform, this guide will help you with principles,

examples, and troubleshooting to effectively prompt AI agents. Note that this guide does not cover AI agent configuration, tool setup, testing, or analytics in great depth. For more information in those areas, view our documentation:

<https://www.servicenow.com/docs/csh?version=latest&topicname=na-ai-agents.html>

Note: The prompt examples provided in this guide, except where noted, do not always constitute complete working instructions for AI agents. They are meant as an excerpt or as an explanation. Many of these examples are based on our out-of-box AI agents. Screenshots in this guide are based on the Yokohama Patch 6 release and may slightly differ than your instance depending on the version.

Acknowledgements

This guide would not be possible without the valuable contributions of the following people at ServiceNow:

- Dan Andrews, Warren Chan, Talia Ben Sasson Gordis, Pradeep Gouribhatla, Mark Griffin, Varun Rapeti, Suhas Sridhar, Alekhya Tethali, TJ Togatapola, and others!

Prerequisites

Before you begin working with Now Assist AI agents, ensure you fulfill the following prerequisites:

- Your ServiceNow instance is on the latest release version that supports AI agents (Xanadu Patch 7 or Yokohama Patch 1 release minimum) with the latest patch and parent Now Assist store apps updated (e.g., Now Assist for ITSM, HRSD...)
 - When discussing AI agent Studio features, this guide assumes you have at least a Xanadu Patch 10/Yokohama Patch 6/Zurich Patch 1 instance version and the latest store app version.
- Your instance user has the `sn_aia_admin` role to access AI agent Studio
- AI Search is enabled
- Now Assist Panel is enabled via the Now Assist admin console
- You have experience with ServiceNow Platform navigation and configuration
- You have a basic understanding of Generative AI, LLMs, and prompt engineering. For an overview, see the ServiceNow University courses on *Introduction to Generative AI* and *Now Assist Essentials* at: <https://learning.servicenow.com/>

Want to learn more about AI agents? Visit our website and schedule a demo at:
<https://www.servicenow.com/products/ai-agents.html>

1. How do Now Assist AI agents work

To effectively prompt agentic workflows and AI agents in the ServiceNow AI Platform, it's essential to understand the core components and how they interact.

An **agentic workflow** contains instructions that determine how a group of connected AI agents can be orchestrated to perform a task. When deciding whether to create an agentic workflow, consider:

- What is the fully defined problem or objective you are trying to solve?
- How would a human or team of humans solve it today?
- What are the success criteria for the problem to be considered solved?

Below is a screenshot that shows an agentic workflow called “Analyze and improve services”.

The screenshot displays the configuration interface for an agentic workflow titled "Analyze and improve services". On the left, a sidebar lists configuration steps: "Define key requirements" (active), "Add a preferred trigger", and "Select a UI display". The main area is titled "Write a clear description and list of steps" and includes a "View versions" dropdown. It contains fields for "Workflow Name" (set to "Analyze and improve services") and "Workflow description" (with a "Description for LLM" label). The description text is: "Analyze customer feedback and derive service metrics, and operational trends to identify improvement opportunities. Deliver actionable recommendations that optimize service delivery and elevate customer satisfaction." Below this is the "List of steps" section, labeled "Version V1" and "Instruction for LLM". It includes a general instruction: "Always ensure you follow this exact same order to deliver the expected result." and two numbered steps: "Step 1: - Engage the Survey Requirement Collector Agent to prompt the user for a start and end date. - Validate and confirm the defined date range for analysis." and "Step 2: - Use the Survey analyzer Agent to extract user survey responses and associated incident details and provide the metrics, trends and actionable recommendations. - Display the output result appropriately". A "Continue" button is at the bottom right.

An **AI agent** is an autonomous, digital worker that carries out instructions using reasoning and a set of defined tools. AI agents contain a **Role**, which defines the job description and personality of the AI agent. The AI agent also contains **Instructions**, which are the steps and information needed for the AI agent to achieve its goal. Below is a screenshot of an AI agent’s Role and Instructions for: “Categorize ITSM incident AI agent”. You can also create and save multiple **versions** of the Role and Instructions as you experiment with the optimal version of your prompts.

Categorize ITSM incident AI agent Chat Exit

Define the role and required steps View versions

Using precise language, define the role for this AI agent and the steps that make up its problem-solving process. For each step, define a specific action and any criteria.

AI agent role * Instruction for LLM

You're an expert in analyzing incidents and are tasked with assigning appropriate category, subcategory and configuration item (CI).

Characters left: 1868

List of steps Version V1 Instruction for LLM

Objective:
Incidents represent disruptions, unavailability of a service, or issues faced by a user. The agent's task is to understand the issue and assign the appropriate:

- Category
- Subcategory
- Configuration Item (CI)

Step 1: Retrieve Incident Details

- Fetch the details of the given incident.
- DO NOT PROCEED if incident details are missing, OR if the incident is in a RESOLVED or CLOSED state.
- If no details are found, verify the correctness of the incident number.
- If the incident is RESOLVED or CLOSED, DO NOT inform the fulfiller and END the flow.
- Only proceed when incident details are successfully retrieved.

Step 2: Assign Category, Subcategory, and CI

Step 2.1: Determine the Category

Continue

Tools are functions and operations that an AI agent can use to achieve its objectives. Tools include descriptions and input descriptions, which serve as prompt instructions. Below is a screenshot of a tool for a Flow Action called "Get details of Change request".

Change implementation plan AI agent Add tool

Add tools and information

Define the role and required steps

Define trigger

Toggle display

Flow Action

Script

Get Similar change requests - Implementation plan

Autonomous false

Fetch change requests which contain similar and relevant information to search query. It should be used only to fetch relevant change requests based on the current change request details. This tool is to be run ONLY ONE TIME to find any relevant change requests related to the issue in the change request.

Get Details of change request

Select flow action *

Get Details of change request

Inputs

Change request Number Data type

change_request_number string

Define these tool settings, which are used only by this AI agent.

Name *

Get details of Change request

Tool description * Description for LLM

This tool fetches the details of the change request given like short description, description, work notes etc.

This tool should only be invoked in the beginning to get details of current change request.

Execution mode *

Supervised Autonomous

Display output *

Yes No

Back Continue

Prompts are processed by an **Orchestrator**, which is a coordinator entity that handles collaboration across multiple agents to accomplish a complex task. The orchestrator employs "Chain of Thought" with the ReAct framework and an LLM to run the AI agents, instructions, and tools. It also provides memory for the team of AI agents to share. You do not need to put system-level prompts in your instructions such as, "You are an agent...think step-by-step...use tools...plan

ahead and reflect..." These basic instructions are already incorporated into the orchestrator.

Note: Due to the non-deterministic nature of LLMs, AI agents may behave differently each time it is run. Consider this when testing AI agents and determine the importance of having either a more static workflow or a dynamic agentic workflow.

Whether to build one large AI agent or several specific AI agents

Depending on your business objectives, you may wonder whether to build one large AI agent with numerous tools, or multiple smaller agents with fewer tools. Think about how you would form a real human team – would you hire one worker who can do many things decently, or multiple workers each with a specialized skillset. You more likely want a team of workers with specialized skillsets. For example, a software development team would have different workers for front-end, back-end, and QA responsibilities. Thus, it's recommended to create smaller agents that can each fulfill one objective very well. AI agent modularity provides you flexibility should you choose to assign the AI agent to different agentic workflows. Smaller AI agents are also helpful for troubleshooting.

From a technical performance standpoint, it is not recommended to have an AI agent with more than 15 tools. At that point, it is better to create separate agents that can complete a specific task.

2. General prompting guidance

Getting prepared

Designing prompts is an iterative process, so it is helpful to have your own outline or design document as a reference. Refer to your prompt drafts and notes in a notepad or in the prompt versions while you build your AI agent. Collaborate with your organization's subject matter experts or business process owners to check whether your AI agent is following the appropriate rules and processes. Consider whether users will trigger and interact with the AI agent, or whether the AI agent will be triggered automatically without user intervention.

Structuring and formatting your prompts

An organized structure to your prompt improves the readability and success rate for an agentic workflow. Use Markdown to label and format your prompt so that it is better understood by human readers and the orchestrator. Markdown contains several syntax elements to help with structuring your prompt:

- Use **#** for Headings, **##** for subheadings, etc. This is recommended to use if your prompt has several sections with navigation instructions to the different sections.
- Important details are noted with ****bold****.
- If you wish to create sections within your prompt, use **---** as a separator.
- Use numbering and sub-numbering such as **1.**, **1.1.**, **2.**, **2.1.**, etc., for any steps the AI agent is expected to follow in sequential order.
- Use dashes, **-**, for bullet points.

Your prompts should provide clear and specific instructions and be written from the 2nd person point of view. Be grammatically correct and avoid overuse of acronyms and abbreviations. Use imperative language, e.g., *"Analyze this"*, instead of *"You should analyze this."* Below is an example of a short, end-to-end prompt where an AI agent first analyzes data from a Universal Request ticket and then generates a resolution plan to the user based on the analysis. The first step, 1.1, uses several tools to complete its objective, such as search retrieval (Knowledge Base search).

Example of an AI agent's instructions:

1. Analysis and Planning

1.1. Initial Analysis

1.1.a. Get details of the incoming universal request ticket.

1.1.b. Extract key information from the request such as short description, description, category, and priority.

1.1.c. Perform Knowledge Base searches for information relevant to addressing the details in the universal request ticket.

1.2. Resolution Planning

1.2.a. Create a Resolution Plan based on the retrieved details for solving the issue in a numbered list format.

1.2.b. Present Resolution Plan to the user for approval.

Note that each step contains as few actions but as much context and descriptive instructions as possible.

Do not leave vague instructions up to the agent's interpretation. For example, don't just say, *"Handle the issue appropriately."* Instead, provide specific instructions like, *"If an error is found, send a notification to the service desk manager."* Using ALL CAPS for certain phrases in the prompt may add weight to the prompt, but it is more helpful to provide specific context and instructions.

Although the agentic workflow and AI agent names and descriptions are not considered prompts, they are still important. Provide a clear and specific name and description for your AI agents. The orchestrator uses this information along with a user's utterance to discover and run an agentic workflow or AI agent in the Now Assist Panel or Now Assist in Virtual Agent. The description should provide a summary of the agentic workflow or AI agent's objectives and actions. Each AI agent should be unique with little to no overlap. Having two AI agents with "task analyzer" in its name with similar instructions, for example, may cause confusion to the orchestrator and the user. Instead, the description should highlight the uniqueness of each AI agent.

Note: It is not necessary to say "please" and "thank you" in your instructions, but it never hurts to be nice to your agent!

This is how a ServiceNow senior AI engineer describes their approach to prompting an AI agent powered by an LLM:

“Imagine the AI agent is a real person, and you are laying out the job description to them along with a set of tools and instructions. The clearer and more concise your instructions, the greater the chance of success for that person. If your instructions confuse your human colleagues, then an LLM would also fall short. The overall goal is to address a non-deterministic problem by providing a structured thought process that can be applied with deterministic tools at scale. If you are getting unexpected results, try splitting up instructions, being concise with simple wording and/or step back and ask whether using a non-deterministic LLM is the best approach to a problem.”

Prompting for agentic workflows

Your agentic workflow instructions should state the objective and how each of the connected AI agents contribute towards the objective. Its instructions do not need to repeat the AI agents' instructions. However, it should provide relevant context to how the AI agents relate to one another. If any repetition is necessary, the language should be consistent. See below for an example of agentic workflow instructions that reference multiple connected AI agents.

Example of agentic workflow's instructions:

*Step 1. Retrieve Record Information***

- *Use the Record Handling Agent to collect all required record details.*
- *Do not proceed until all necessary task information is successfully retrieved.*

*Step 2. Document Resolution Steps***

- *Use the Next Best Action Agent to generate the complete resolution steps for the issue.*
- *Clearly display the generated resolution steps to the user in a bulleted list format.*
- *If resolution steps are unavailable, update "work_notes" for the given record number and table name, with the following message: "I couldn't generate a summary due to a lack of required resources."*

...

Prompting for AI agents

The AI agent's **Role** gives the agent a persona which allows for more context-aware responses. For example, you can say, “You are an IT security expert...who specializes in...”, and the agent and orchestrator will understand the domain and level of expertise desired by the user. It is not necessary to provide action steps in

the Role, but it is a good place to provide background information, target audience, and guardrails. Below is an example with optional sections.

Example of AI agent role:

You are an expert ServiceNow Self-Service Resolution Agent specializing in creating user-friendly instructions and managing self-service resolution processes. You function as an empathetic technical translator who converts complex solutions into simple language that all users can understand.

(Optional)

Core Functions – You are proficient at:

...

Constraints – You are NOT proficient at:

...

AI agent instructions should state the objective to be completed and tools to be used. It also should contain as much information, context, and guardrails needed by an AI agent to effectively complete its objective. Tool names and variables that are expected to be used in a step should be called out.

The following is an example template plus language of different sections to provide your instructions depending on the complexity of your agent. Try not to combine or mix sections to avoid confusion by the orchestrator or the prompt writer:

Objectives

“Your objective is to...”, “You will help with...by doing ...”, “You should generate...”

Validations [prior to execution]

“First check if...”

Steps

- “1. Use the “Search Knowledge” tool to look up...”*
- “2. If no articles are found, proceed to update...”*

This will be the main body of the prompt and is the minimally required information for an AI agent to execute successfully. See more guidance regarding Steps below.

Expected Output

“Present a resolution plan with the following information to the user...”, “The output format should be...”

Validations for after execution, i.e., Success Criteria

"Before finalizing/ending, verify that..."

Constraints

Other Do's and Don'ts not covered in the Steps.

Additional instructions and constraints

(See the chapter, "More prompt examples", to see prompts that take cues from this structure.)

Each step should clearly describe what needs to be done and any additional context and logic. If a step needs to be fulfilled by a specific Tool, the tool name and action should be specified. If a step's intent depends on a previously completed step, clearly define that at the beginning of the step. This acts as a "gate" for the following steps. Clearly state if a step outcome determines or influences anything in the following steps, or if the Agent should not continue until a particular action or sub-objective is achieved. You can use conditional "If/Then" throughout your instructions to provide different actions to your AI agent depending on responses and situations. Be careful with combining multiple actions in a single step or in a run-on sentence. This may confuse the orchestrator, and important actions may be skipped.

The following is an example of a prompt that references several Tools, e.g., "Get similar tasks" in step 2, that also provides additional actions and logic:

Example of AI agent instructions:

- 1. Get record details for the given task. If the task details are not found, then check the correctness of the task number from user. Do NOT move forward with the steps until details of task are collected.*
- 2. Next, fetch a list of similar records to understand how they were resolved.*
- 3. Next, fetch a list of relevant knowledge articles, based on the short description of the task. IMPORTANT: You have enough information to create a resolution summary steps if there are either similar tasks or similar knowledge articles or both.*
- 4. If details of both similar tasks AND similar knowledge articles are NOT present, show the output "I couldn't come up with a plan as I didn't have the required resources." to the user and then end.*

(Spelling out the exact tool name is not always necessary as seen here, but it is otherwise helpful for the orchestrator and the prompt writer.) See the chapter “More prompt examples” for a longer version of this prompt.

Finally, to end an AI agent execution, simply write a step that says the agent should end. Though explicitly stating an end is not required, it is recommended so that there is no ambiguity, and the orchestrator does not try to repeat or prolong an AI agent execution.

Other prompting techniques

The following are specific techniques that will be useful to you when writing prompts for your AI agents to fulfill certain objectives.

Referencing ServiceNow data

If your AI agent is interacting with ServiceNow platform data, your prompt can refer to ServiceNow tables and fields as they are recognized by the orchestrator. This is effective in combination with tools such as Script or Record Operation tools. That said, an AI agent cannot look up the data itself; it requires a tool. See the following example that retrieves data from an Incident record provided by a user query via a “Get Incidents” flow action tool:

Example of AI agent instructions referencing a tool and ServiceNow data:

Step 1: Get the Incident record

Step 2: Review the Incident description for any misspellings or grammatical errors.

In some cases, it is possible to refer to specific record sys_IDs, however that is not scalable (the AI agent may need to look at different records in the future), so it is not a recommended practice. See the “More prompt examples” chapter for another example of how the orchestrator can retrieve and use data it retrieves from tools.

Handing over from one agent to another

Within an agentic workflow, it is possible to pass the user between different agents, so that one agent does not do all the work. This can be done in the agentic workflow instructions or the AI agent instructions. See the example below from an Agentic Workflow with two AI agents that retrieves record information

and generates resolution steps respectively. The AI agent names have been **emphasized**.

Example of agentic workflow instructions for multiple AI agents:

Step 1. Retrieve Record Information

- Use the **Record Handling Agent** to collect all required record details.
- Do not proceed until all necessary task information is successfully retrieved.

Step 2. Obtain Resolution Steps

- Use the **Next Best Action Agent** to generate the complete resolution steps for the issue.
- Clearly display the generated resolution steps to the user

Using memory

Memory is used to store relevant information throughout the Agentic workflow. After you run a tool that retrieves data such as "Get Incident Description", you can reference that data later in your instructions, such as, "Rewrite the Incident description in the following manner..." The orchestrator is aware of the 'Incident description' object from the earlier instructions and tools.

If you're juggling multiple data references, be clear as to what data you're referring to when prompting, whether it's a User ID or an Asset description.

Example of memory in prompts

Step 1: Rewrite the Incident description in the voice of a science professor

...

Step 5: Update the Incident description with the new information.

Memory can be referenced across different AI agents in an agentic workflow. The AI Agent Studio can also store long-term memory in the ServiceNow instance. This is configurable in the Settings page in "AI agent data access".

End-user inputs and outputs

Often times, the AI agent will need input from the user. The orchestrator will recognize the term "user" in the prompt as the end-user. It is not recommended to refer to the user by their title or role in the prompt, e.g., "Ask the Service Desk

Agent Tier II...", as it may confuse the orchestrator whether to refer to the end-user or another AI agent etc. Below are example prompts to ask the end-user for inputs such as data or approval.

Examples (with notes):

"Ask the user to provide the Incident number..."

(You would then likely instruct the AI agent to pass that data to a tool, to another step, or another AI agent.)

"Ask the user for their approval (or feedback) before continuing to the next step...Refine output with user guidance...until the user is satisfied."

(The orchestrator is smart enough to discern the sentiment of the user's response in determining the approval, i.e., the orchestrator understands responses like "yes" or "no". If more specific approvals or feedback are required from the end-user, then provide that instruction in the prompts. You can use this pattern to let users ask the AI agent to refine responses with each feedback.)

In your prompt, you can create a menu of more than one action a user can take. Each action option should then have a section of instructions and possibly tools. See the example below where the AI agent asks whether a user wants to add a comment to a ticket (via a tool) or skip that action.

Example:

... 2.6 Ask from user the action the user wants to perform on the ticket after the ticket details. Always have "Add a Comment" as action. Always allow the user to skip if no action is desired.

3. Add a Comment to a Ticket

3.1. Ask for the comment from the user if not provided.

3.2. Mandatory confirmation:

- After executing the Add a Comment action, you must always inform the user of the outcome.*
- Clearly state to the user that the comment addition has been completed (or failed) based on the tool output.*

To display information to the end-user from the AI agent, you can use a prompt like, *"If an Incident record is not found, inform the user (or, output to the user) that the Incident record could not be found."* If you need the output to be specific, use quotation marks. You can also rely on the orchestrator LLM to generate content to output to the user, e.g., *"Create a professional summary of the Incident details to show the user. It should be in the following format..."* If the output is coming from a tool, reference the tool name in the instructions. You can provide instructions on the length and tone of the output. You can also provide specific and polished outputs via formatting instructions such as the examples below. Note that the AI agent orchestrator can currently only input and output string.

Example of a specific format in instructions:

Create a professional yet conversational completion summary with this format to show the user:

TASK CREATION STATUS REPORT

****Original Request**:** [Original Ticket Number]

****New Case**:** [Incident/HR Case #] [New Case Number]

****Category**:** [Category Used] (in italics)

Another example of HTML formatting instructions:

Merge the selected steps into a single, ordered, non-redundant list, in the following format and save to memory:

Recommended Resolution Plan:

Step 1

Step 2



Step 3

...

Using Now Assist to create a prompt

You can use Now Assist to get a starting point in writing a prompt for your AI agents. When you create a new AI agent, you can click the button, **"Generate description, role, and steps"** and then provide a general description of what you

want the AI agent to do. Then click **“Generate”** and the Description, Role, and Instructions will be automatically generated. You can then expand and adjust the prompt (perhaps using this guide!) to optimize its accuracy and success.



3. Prompting with Tools

Prompt instructions are used in the **Description** field of almost every tool. The Tool description instructs when and why a tool is used by an AI agent. It also instructs the AI agent on how to use its inputs and outputs. Many of the prompt guidance for AI agents also apply to Tools. The tools and instructions mapped to the AI agent are automatically visible to the orchestrator, so avoid redundant prompts like “You are given these tools” or “You are provided with the below tools”. You also do not need to repeat instructions in both the AI agent instructions and in the Tool description. Like AI agents, the name of the Tool is relevant; do not give it a generic name like, “Tool 1”. Avoid overly similar tools names and descriptions. Creating more descriptive tool names lets the orchestrator better align the AI agent instructions with the Tool to be used without having to be overly explicit every time. Below are examples of tool names and instructions, such as tools that get similar records via a flow action from an input record number.

Examples (with notes):

Tool Name: Get Similar Tasks

Tool description: This tool is used to find the Tasks that are similar to the given type of task. You will be passed table and record number as inputs.

Tool Name: Course Cost eligibility check

Tool Description: This tool verifies if a course cost qualifies under a given requester's location's maximum annual benefit. This tool should only be called if the checklist includes maximum annual benefit criterion.

(For a course tuition reimbursement agentic workflow.)

Tool Name: Create Change request record

Tool Description: Create a change request record with the resolution plan

(The resolution plan was generated from a previously run tool and serves as an input for this tool to create change request records. Below is a screenshot of the example Tool description and inputs.)

Note: One of the tool input's descriptions seen below (*change_title*) instructs the orchestrator to generate a summary of the resolution plan. A

separate tool is not necessary for summarization. The orchestrator can directly perform simple LLM operations like summarizing or determining positive/negative sentiment. If you require a more complex prompt to generate an output without weighing down a tool or AI agent prompt, consider "outsourcing" it to a custom Now Assist skill built via Skill Kit. That skill can then be used as an AI agent tool.

The screenshot shows a form titled "Add a record operation" with the subtitle "Fundamental actions that allow AI agents to interact with and manage data." The form contains the following fields:

- Name ***: A text input field containing "Create Change Request record".
- Description ***: A text input field containing "Create a change request record with the resolution plan". To the right of this field is a blue button labeled "AI instruction".
- Inputs**: A section containing two input pairs:
 - Input name ***: "change_title" and **Description**: "Summary of resolution plan".
 - Input name ***: "res_plan" and **Description**: "Resolution plan created by Next Best Action Agent".

More tool description examples are available in the "More prompt examples" chapter.

When your tool contains inputs like above, instruct how the inputs are used by the tool in its Description field. Describing the purposes of the inputs strengthens the orchestrator's understanding of how to use the tools. When defining an input name, give them unique names, even if same as the field name, e.g., "incident_number". That way, you provide enough context to the orchestrator so that they can be differentiated during processing. If an input or output is mandatory, state so in the Agent instructions or Tool description.

For some tools, it is recommended to inform the user of the outcome such as a success or failure. That way the user is aware of the AI agent's progress. You can do so with a prompt in the AI agent like, "After executing the <...> tool, you must always inform the user of the outcome."

Where possible, use tools to do data filtering, look-up, or manipulations, instead of relying on a prompt and the orchestrator. For example, if you want to find all incidents assigned to a certain user, use a tool instead of trying to prompt instructions on how an AI agent should find and filter that data.

List of internal AI agent tools

The AI agent orchestrator utilizes internal tools in the background to process AI agent instructions and provide behind-the-scenes support for performance and user experience. You will see these tools mentioned in the logs during AI agent test runs. Do not name your tools like these internal tools listed below as it may cause confusion with the orchestrator.

- *Organize_general_knowledge*
- *Math*
- *Fallback*
- *Finish*
- *Join*
- *Generate_content*
- *Check_with_other_agents*
- *Communicator_agent*
- Other problematic names for your tools: *Content analysis*, *User input*, *User output*

Note with caution: Some of these internal tools allow you to query the LLM's general knowledge with your AI agent's instructions without having to use a tool. This may be helpful to retrieve information that is not found in your ServiceNow instance, but it is also more open to non-relevant results.

List of keyword verbs to use for tools

Below is a sample of verbs to use in your instructions to prompt certain kinds of tools you will use in your AI agents.

Keyword verbs	Actions
<i>analyze, evaluate, assess</i>	Triggers analytical tools
<i>fetch, retrieve, filter</i>	Activates data tools
<i>generate, synthesize, compile</i>	Enables creation tools
<i>verify, validate, confirm</i>	Invokes validation tools
<i>show, display, gather</i>	Triggers UI tools

4. More prompt examples

Now that you have a solid foundation for prompting and troubleshooting, let's review more prompt examples. ServiceNow also provides dozens of out-of-box agentic workflows and AI agents across many workflows (IT Service Management, HR Service Delivery, CRM) and industries. Use these AI agents as a starting point or as inspiration for your own agentic workflows. You can also duplicate these out-of-box AI agents to make them your own.

Review the following prompt draft. See if you find any issues with it. There's an issue in almost every step:

Example draft instructions with room for improvement:

To check and allocate a software license, follow these steps:

Step 1: If there's license id, then proceed to step 2.

- If not provided, then end.

Step 2: Follow the following steps:

- If this is manager user, proceed to step 5.

- If this is regular user, proceed to step 4.

- If this is guest user, proceed to step 2.

Step 3: Guest user license allocation:

- Use an appropriate tool to check if the user license can be allocated to a guest user

- If guest user can be allocated, tell user that you are assigning it to a guest user.

- Proceed to step 6.

Step 4: Manager user license allocation:

- If the manager is allowed to get a license, check if a license exists.

- If it does, assign it and let the user know.

- Otherwise, do something else.

- Move to Step 6.

Step 5: Regular User License Allocation:

- Check if the requested software license is there.

- If available, the user should get it.

- If not, then go to Step 6.

Step 6: Handle Something (Unclear Purpose)

- If everything is fine, then move forward.

- If something went wrong, figure out what it is and fix it.

Did you catch all the issues? Here are just some of them:

- In Step 1, how would the AI agent know whether a user already has a license ID or not or how to find a license ID? This needs instructions for user input or a tool to find this information. It would also be helpful to describe the properties of a License ID for the orchestrator or AI agent to identify. Don't make assumptions.
- In Step 2, instructions like "follow the following steps" are unnecessary. Also, the last bullet contains a typo: it incorrectly loops into Step 2 again.
- In Step 3, it says "use the appropriate tool" which is vague. It would be better to call out the tool name or just say, "Check if the user license can be...", and make sure there is a tool name aligned with this instruction.
- In Step 4, the instructions are missing the criteria for whether a manager is "allowed" to get a license. Otherwise, this needs a tool again. Also, "do something else" is too vague. Make sure your instructions do not have leftover notes or placeholder text.
- In Step 5, there is more vagueness with "software license is there" and "user should get it." Table names and tool operations should be stated more thoroughly.
- Step 6 is unnecessary and incorrectly used as a catch-all with unclear instructions.

See below for a revised version of the above steps that is more understandable by both the orchestrator and a prompt writer:

Example revised instructions:

To verify the eligibility criteria and allocate a software license, follow the below instructions step-by-step:

Step 1: Verify if the License ID is available on the user record. If the License ID is missing, ask the user to provide it.

Step 2: Fetch the user details to determine the user type.

Step 3: If the user type is 'manager', then follow the below steps:

Step 3.1: Check if the requested software license is available.

Step 3.2: If available, allocate it to the user.

Step 3.2.1: Once the license is successfully allocated, inform the user about the same.

Step 3.3: If the licenses are not available, then inform the user the requested licenses are not available.

Step 4: If the user type is 'regular employee', then follow the below steps:

Step 4.1: Create a license allotment task for the user.

Step 4.2: Once the task is created, inform the user that the license allotment task has been created, and it is waiting for their manager's approval. The license allocation will be processed based on their manager's approval.

Step 5: If the user type is 'guest', then follow the below steps:

Step 5.1: Verify the roles of the guest user.

Step 5.2: Retrieve the list of roles who are eligible for the requested software license.

Step 5.3: Compare the guest user's roles with the software license eligibility roles to determine if the guest user is eligible for the requested license.

Step 5.4: If the guest user is eligible for the requested license, and then create a temporary license allocation task.

Step 5.4.1: Inform the user that the temporary license allocation task has been created for the guest user.

Step 5.5: If the guest user is not eligible for the requested license, then inform the user that the guest user is not qualified to access the requested software license due to missing roles.

The following is another example of a complete AI agent prompt for resolving tasks. A briefer version was shown in the chapter "General prompting guidance". The prompt starts with context: defining what a Task is and the overall objective (some of this could also be in the description or an Agentic workflow prompt). The prompt is organized in sections and has several steps with contextual information and directions on what to do when certain criteria are met or not met, i.e., "gates". It accounts for situations where criteria are not met and provides a fallback action. It also provides additional instructions.

Example:

Tasks are records for tracking issues and their resolution updates through a set of journal entries. Your objective is to understand what the issue with the incoming task is, search for potential solutions and come up with a good plan for solving the issue presented in the task.

#Steps:

To resolve a task, ALWAYS follow this progression of steps:

1. Get record details for the given task. If the task details are not found, then check the correctness of the task number from user. Do NOT move forward with the steps until details of task are collected.

2. Next, fetch a list of similar tasks. If similar tasks are present, understand the general procedure followed ONLY as a reference.

3. Next, fetch a list of relevant knowledge articles, based on the short description of the task. If similar knowledge articles are present, understand the general procedure followed ONLY as a reference.

IMPORTANT: After gathering the above details, make sure to check if you have enough information to create a plan. You have enough information if there are similar tasks or similar knowledge articles.

If similar tasks OR relevant knowledge articles are present:

1. MANDATORY: Now, based on all of the above information, generate a plan for resolving the task in a numbered list format.

2. Inform the user of the plan and collect feedback or approval. Make sure that you explicitly mention in the Thought that you need to reach out the user. ALWAYS treat user's response to this as a critique to the current plan and NOT as permission or direction to proceed.

3. Based on the user feedback, revise the plan and generate a new one in a numbered list.

4. Inform the user of the revised plan and collect feedback or approval. Make sure that you explicitly mention in the Thought that you need to reach out the user. ALWAYS treat user's response to this as a critique to the current plan and NOT as permission or direction to proceed.

5. Repeat step 3 and 4 until the user approves the plan. Keep revising the plan and keep collecting user feedback until the user approves the plan. Do NOT proceed with the plan until the user confirms.

6. ONLY after the user has approved the plan, Finish the execution with an Observation as "With all the available resources, I have shown the plan. Let me know what else I can help you with."

If similar tasks AND similar knowledge articles are NOT present,

7. Finish the execution with an Observation as "I couldn't come up with a plan as I didn't have the required resources. Let me know what else I can help you with."

#Additional instructions:

- ALWAYS reach out to the user to confirm the plan before ending the execution.

- NEVER end execution until the user explicitly confirms that the plan is good to go. ALWAYS ask the user if they're satisfied with the plan.

You'll have at least three tools for this prompt: getting the details of the task, getting similar tasks, and fetching relevant knowledge articles.

The following is an example of a Task Creation Agent, which takes as input a Universal Request (UR), retrieves its field data, uses the orchestrator to categorize it as either an "IT Incident" or "HR Case", and then creates either an Incident or HR Case based on the results.

Example:

Step 1: Analyze Universal Request

- Retrieve the Universal Request details - short description and description and determine whether the request should lead to creating an HR Case or an IT Incident.

- Example descriptions for HR Case: "paycheck", "leave policy", "manager", "performance review". If the Universal Request fits this description, categorize the Universal Request as "HR" ticket type.

- Example descriptions for IT Incident: "laptop", "email", "Microsoft". If the Universal Request fits this description, categorize the Universal Request as "Incident" ticket type.

Step 2: Use the "Create Incident or HR Case" Tool

- Prepare the following parameters from the Universal Request:

** Ticket Type: Use either "Incident" or "HR" as determined in Step 1*

** Work Notes: Generate work notes that include:*

- Brief summary of the original request

- Justification for creating this specific ticket type

- Key findings from the analysis

- Any special handling instructions

- Call the "Create Incident or HR Case" tool with these parameters

- Store the new ticket/case number for reference in subsequent steps

Step 3: Verify and Close Process

- Verify ticket creation was successful by confirming:

** New ticket/case number was received*

** All information from the Universal Request was properly transferred*

** No errors were returned by the "Create Incident or HR Case" tool*

- Create a professional yet conversational completion summary to show the user the steps taken.

#Additional Instructions

- If any errors occurred during ticket creation, tell the user: "I encountered an issue while trying to create the new ticket based on your Universal Request."

There are two tools used by this AI agent: "Get Universal Request Details" and "Create Incident or HR Case". Based on the data retrieved from the Universal Request, the instructions then ask the Orchestrator to categorize the UR as either "HR" or "IT" ticket type based on the UR's content. The instructions provide example descriptions as guidance. It also provides validation gates and fallback messages. The inputs for the Create Incident or HR Case tool (a Subflow) would be the UR details and the orchestrator-determined ticket type, "HR" or "IT". It then uses these inputs, and an IF/THEN control in the Subflow to either create an Incident or an HR Case, generates a Work Note to that new ticket, and then outputs the results details to the user. The tool's instructions would be something like the following:

Example:

Tool name: Create Incident or HR Case

Tool description: Use this tool to create either an Incident or an HR Case.

Understanding tool inputs:

Ticket Number: The Universal Request number you are working on

Ticket Type: This will be either HR or Incident

Work Notes: Provide a summary of the Universal Request you are working on and a quick summary of why you are creating the new Incident or HR Case. This will help the human reviewing the new ticket understand why this was created.

##MUST DO:

When creating work notes for the new ticket, you must always begin your notes with your name: "Task Creation Agent:" followed by the actual note content."

5. Prompting troubleshooting

Don't be dissuaded if your AI agent doesn't behave as expected. Building, testing, and deploying agents is an iterative process. When testing your agentic workflow or AI agent in the Studio, you can use the Decision Log on the right-hand side to see the behind-the-scenes thinking of your AI agent. You can also navigate to the "Activity" tab in AI Agent Studio to review and replay past agentic executions.

If there's an issue, the first thing to always check is that you are on the latest available store application version of *Now Assist AI Agents* for your instance version. If an incorrect agent or no agent is found during a run in Now Assist Panel or Virtual Agent, check the uniqueness of the agent's names and descriptions. Don't forget that the AI agent Name and Description is used by the orchestrator in identifying the right agent for the right task. If there are multiple agents with overlapping or unclear context, that might confuse the orchestrator regarding which agent is most suitable for the task. Watch out for misspellings or unintentional formatting errors that would cause confusion for an orchestrator just as it would for a human reader.

If an incorrect tool or no tool is used during a run, check the tool name and description. Check that the agent prompt and tool description properly reference the inputs and outputs and how they're used. Make sure there is no confusion with similarly named tools or variables. If your tool appears stuck, you may be processing too much data. Try limiting the input data amount and number of returned results via conditions and filters. Consider adding fallbacks as instruction steps to catch when a tool errors out or returns no results.

If your output is poorly formatted, check the tool's Output transformation strategy field. That field allows you to provide tool outputs in Concise, Verbose, or as given (None) format. Sometimes, the 'None' option defaults to JSON-formatted output, so use Concise or Verbose. Otherwise, your prompt should instruct how to format

the output, e.g., *“Output all retrieved information in a well-formatted, numbered message.”*

If your AI agent is looping, make sure your prompt is not too open-ended or without a concluding action. It may help to add a gate in the instructions or an explicit “end” step.

If your agent's tests are still inconsistent, you may need to provide some repetition or more explicit instructions. For example, we note earlier that you don't need to say, “use the available tools... to get the user details...”, but it may help to put in some of this language if the tool is running inconsistently. It may also help to reiterate objectives, roles, and important guardrails at the beginning of a prompt prior to the list of steps. A reminder: LLMs are non-deterministic, so the more complex the AI agent or the more tools added to the AI agent or agentic workflow, the more unpredictable the behavior may become. But with a clear and concise voice in your AI agent instructions, you can create powerful and dynamic agentic workflows that achieve your business objectives at massive scale.

6. Additional Resources

To learn more about AI agents, see the resources available to further your understanding of how to build and deploy AI agents.

- Official Documentation for ServiceNow AI agents:
<https://www.servicenow.com/docs/csh?version=latest&topicname=na-ai-agents.html>
- ServiceNow University: <https://learning.servicenow.com/>
- Now Assist Community Forum:
<https://www.servicenow.com/community/now-assist/ct-p/now-assist>

Excited to unlock productivity at scale with AI agents? Learn more at:

<https://www.servicenow.com/products/ai-agents.html>

Schedule a demo with us at: <https://www.servicenow.com/contact-us-generative-ai.html>

